

```

In [ ]: #IMPORTING BANK DATASET

import pandas as pd
bank = pd.read_csv(r'C:\Users\NJERI SHAWN\Desktop\python pdf\Bank.csv')
print(bank)
bank['Salary'].mean() #same as sal.mean() after assigning
sal=bank['salary']
sal=bank['Salary']
sal.mean()
sal.min()
sal.max()
sal.median()
sal.std()

#or we can get all statistical summary of all numeric data in the dataset
by using the describe() command
bank.describe()

```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	\
0	1	3	1	92	69	Male	1	No	
1	2	1	1	81	57	Female	1	No	
2	3	1	1	83	60	Female	0	No	
3	4	2	1	87	55	Female	7	No	
4	5	3	1	92	67	Male	0	No	
..	...	...	...	...	...	...	...	...	
203	204	3	6	61	35	Male	0	No	
204	205	5	6	59	34	Male	0	No	
205	206	5	6	63	33	Male	0	No	
206	207	5	6	60	36	Male	0	No	
207	208	5	6	62	33	Female	0	No	

	Salary	Mgmt
0	32.0	Non-Mgmt
1	39.1	Non-Mgmt
2	33.2	Non-Mgmt
3	30.6	Non-Mgmt
4	29.0	Non-Mgmt
..	...	...
203	95.0	Mgmt
204	97.0	Mgmt
205	88.0	Mgmt
206	94.0	Mgmt
207	30.0	Mgmt

[208 rows x 10 columns]

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	YrsPrior	Salary
<b>count</b>	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000	208.000000
<b>mean</b>	104.500000	3.158654	2.759615	85.326923	54.605769	2.375000	39.921923
<b>std</b>	60.188592	1.467464	1.566529	6.987832	10.318988	3.135237	11.256154
<b>min</b>	1.000000	1.000000	1.000000	56.000000	30.000000	0.000000	26.700000
<b>25%</b>	52.750000	2.000000	1.000000	82.000000	47.750000	0.000000	33.000000
<b>50%</b>	104.500000	3.000000	3.000000	87.000000	56.500000	1.000000	37.000000
<b>75%</b>	156.250000	5.000000	4.000000	90.000000	63.000000	4.000000	44.000000
<b>max</b>	208.000000	5.000000	6.000000	93.000000	73.000000	18.000000	97.000000

```
In [ ]: sal.mean(),sal.min(),sal.max(), sal.std()
```

```
Out[ ]: (39.92192307692307, 26.7, 97.0, 11.256153921958742)
```

```
In [ ]: sal.describe()
```


```
Out[ ]: count      208.000000
mean         39.921923
std          11.256154
min          26.700000
25%          33.000000
50%          37.000000
75%          44.000000
max          97.000000
Name: Salary, dtype: float64
```

```
In [ ]: bank['Salary'].describe()
```


```
Out[ ]: count      208.000000
mean         39.921923
std          11.256154
min          26.700000
25%          33.000000
50%          37.000000
75%          44.000000
max          97.000000
Name: Salary, dtype: float64
```

```
In [ ]: import seaborn as sns
sns.histplot(x=bank['Salary'])
```


```
Out[ ]: <Axes: xlabel='Salary', ylabel='Count'>
```




```
In [ ]: sns.histplot(x=bank['Salary'], bins=10, kde=True);
```




```
In [ ]: sns.histplot(x=bank['Salary'], bins=10, kde=True, color='brown');
```



```
In [ ]: sns.boxplot(y=bank['Salary']);
```



```
In [ ]: sns.boxplot(y=bank['Salary'], color='lightgreen', showmeans=True);
```



```
In [ ]: #Filtering Data in Python
```

```
In [ ]: bank['Gender']=='Female'
```

```
Out[ ]: 0    False
1     True
2     True
3     True
4    False
...
203   False
204   False
205   False
206   False
207    True
Name: Gender, Length: 208, dtype: bool
```

```
In [ ]: bank[bank['Gender']=='Female']
```

Out[ ]:

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt
<b>1</b>	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt
<b>2</b>	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt
<b>3</b>	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt
<b>5</b>	6	3	1	92	71	Female	0	No	30.5	Non-Mgmt
<b>6</b>	7	3	1	91	68	Female	0	No	30.0	Non-Mgmt
...	...	...	...	...	...	...	...	...	...	...
<b>186</b>	187	5	5	86	58	Female	2	No	50.0	Mgmt
<b>187</b>	188	5	5	83	49	Female	2	No	61.8	Mgmt
<b>188</b>	189	4	5	79	52	Female	0	No	43.0	Mgmt
<b>190</b>	191	5	5	86	58	Female	6	No	58.5	Mgmt
<b>207</b>	208	5	6	62	33	Female	0	No	30.0	Mgmt

140 rows × 10 columns

In [ ]:

```
FemaleEmployees=bank[bank['Gender']=="Female"]
type(FemaleEmployees)
```

Out[ ]:

```
pandas.core.frame.DataFrame
```

In [ ]:

```
round(FemaleEmployees['Salary'].mean(),2)
```

Out[ ]:

```
37.21
```

In [ ]:

```
(bank['Gender']=='Female')&(bank['JobGrade']==1)
```

Out[ ]:

```
0    False
1     True
2     True
3     True
4    False
...
203  False
204  False
205  False
206  False
207  False
Length: 208, dtype: bool
```

```
In [ ]: bank[(bank['Gender']=='Female')&(bank['JobGrade']=='1')].shape
```

```
Out[ ]: (48, 10)
```

```
In [ ]: bank[bank['JobGrade']>=4]
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt
<b>145</b>	146	5	4	90	62	Male	3	No	44.5	Non-Mgmt
<b>146</b>	147	5	4	91	65	Male	1	No	41.0	Non-Mgmt
<b>147</b>	148	5	4	89	58	Male	3	No	44.0	Non-Mgmt
<b>148</b>	149	5	4	89	65	Male	0	No	44.0	Non-Mgmt
<b>149</b>	150	5	4	90	63	Female	4	No	42.5	Non-Mgmt
...	...	...	...	...	...	...	...	...	...	...
<b>203</b>	204	3	6	61	35	Male	0	No	95.0	Mgmt
<b>204</b>	205	5	6	59	34	Male	0	No	97.0	Mgmt
<b>205</b>	206	5	6	63	33	Male	0	No	88.0	Mgmt
<b>206</b>	207	5	6	60	36	Male	0	No	94.0	Mgmt
<b>207</b>	208	5	6	62	33	Female	0	No	30.0	Mgmt

63 rows × 10 columns

```
In [ ]: mgmt=[4,5,6]
bank[bank['JobGrade'].isin(mgmt)]

#isin() method checks if the dataframe contains the specified values
```

Out[ ]:

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt
<b>145</b>	146	5	4	90	62	Male	3	No	44.5	Non-Mgmt
<b>146</b>	147	5	4	91	65	Male	1	No	41.0	Non-Mgmt
<b>147</b>	148	5	4	89	58	Male	3	No	44.0	Non-Mgmt
<b>148</b>	149	5	4	89	65	Male	0	No	44.0	Non-Mgmt
<b>149</b>	150	5	4	90	63	Female	4	No	42.5	Non-Mgmt
...	...	...	...	...	...	...	...	...	...	...
<b>203</b>	204	3	6	61	35	Male	0	No	95.0	Mgmt
<b>204</b>	205	5	6	59	34	Male	0	No	97.0	Mgmt
<b>205</b>	206	5	6	63	33	Male	0	No	88.0	Mgmt
<b>206</b>	207	5	6	60	36	Male	0	No	94.0	Mgmt
<b>207</b>	208	5	6	62	33	Female	0	No	30.0	Mgmt

63 rows × 10 columns

In [ ]:

```
#Recoding Data in Python
```

In [ ]:

```
#Adding a Column in our Dataset
bank['Dummy']=0
bank.head()
```

Out[ ]:

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	Dummy
<b>0</b>	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	0
<b>1</b>	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	0
<b>2</b>	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	0
<b>3</b>	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	0
<b>4</b>	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	0

In [ ]:

```
#Dropping a Column in our Dataset
bank.drop('Dummy',axis=1, inplace=True)
```



```
bank.head()
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt

```
In [ ]: #Recoding using the numpy where method
import numpy as np
bank['GenderDummy_F'] = np.where(bank['Gender']=="Female",1,0)
bank.head()
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	GenderD
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	

```
In [ ]: #Recoding Using the apply() Function
#The easiest way to see how this works is to start with a parameterized
function that implements the if/then logic. What follows is a standard
function declaration in Python. The code defines a new function
called "my_recode" which takes a single parameter "gender". The function
returns a 1 or 0 depending on the value passed to it:

def my_recode(gender):
    if gender == "Female":
        return 1
    else:
        return 0
```

```
In [ ]: my_recode("Female"), my_recode("Male")
```

```
Out[ ]: (1, 0)
```

```
In [ ]: bank['GenderDummy_F']=bank['Gender'].apply(my_recode)
bank.head()
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	GenderD
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	

```
In [ ]: #Recoding Using a Lambda Function
bank['GenderDummy_F']=bank['Gender'].apply(lambda x: 1 if x == "Female"
else 0)
bank.head()
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	GenderD
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	

```
In [ ]: #Replacing Values from a List
grades=[1,2,3,4,5,6]
status=["non-mgmt","non-mgmt","non-mgmt","non-mgmt","mgmt","mgmt"]
```

```
bank['Manager']=bank['JobGrade'].replace(grades,status)
bank.head()
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	GenderD
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	

```
In [ ]:
```

```
genders=["Female", "Male"]
dummy_vars=[1,0]
bank['GenderDummy_F'] = bank['Gender'].replace(genders, dummy_vars)
bank.head()
```

```
Out[ ]:
```

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	GenderD
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	

```
In [ ]:
```


```
#Logging variables
bank['logSalary']=np.log(bank['Salary'])
bank.head()
```

Out[ ]:

	Employee	EducLev	JobGrade	YrHired	YrBorn	Gender	YrsPrior	PCJob	Salary	Mgmt	GenderD
0	1	3	1	92	69	Male	1	No	32.0	Non-Mgmt	
1	2	1	1	81	57	Female	1	No	39.1	Non-Mgmt	
2	3	1	1	83	60	Female	0	No	33.2	Non-Mgmt	
3	4	2	1	87	55	Female	7	No	30.6	Non-Mgmt	
4	5	3	1	92	67	Male	0	No	29.0	Non-Mgmt	

In [ ]:

```
import seaborn as sns
sns.kdeplot(x=bank['logSalary'], fill=True, linewidth=2);
```



## Gap analysis with Continuous Variables

In [ ]:

```
# Recall the purpose of Gap Analysis: determine whether two samples of data are different. In our running example, we want to determine whether Sample 1 (salaries of female employees in the bank) is different from Sample 2 (salaries of male employees at the bank). We generally come at
```

Gap Analysis in two steps:

# 1. Plot the data in such a way that we can visually assess whether a gap exists. These visualizations also come in handy later when communicating the results of any formal analysis.


# 2. Conduct a formal gap analysis using statistical techniques

#Using BoxPlots

#ensure Seaborn is loaded

```
import seaborn as sns
```


```
sns.boxplot(x=bank['Salary'], y=bank['Gender'], showmeans=True);
```



In [ ]:

#As an aside: We can do the same kind of analysis by "JobGrade". But recall that we left JobGrade as an integer and did not convert it to a category variable (as we did for Gender and PCJob). We can make this conversion on the fly in order to get a boxplot:


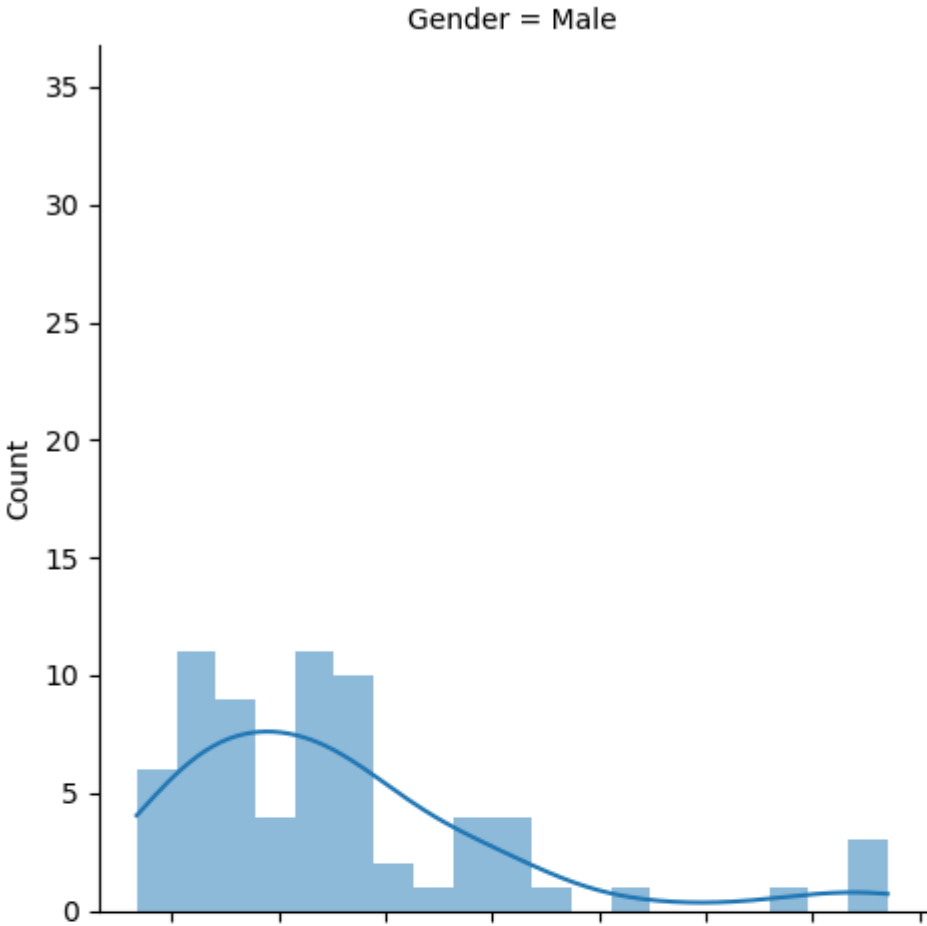
```
sns.boxplot(x=bank['Salary'], y=bank['JobGrade'].astype('category'), showmeans=True);
```



In [ ]:


```
# Faceted histograms
# We used the notion of a "facet" in R to create a grid of histograms. In
this case, we want a grid with one column and two rows. The rows correspond
to different values of the "Gender" variable. This is a bit easier in
Python with Seaborn's displot function, which creates a faceted
distribution plot. Here the row argument tells Seaborn to create one row
foreach value of gender. I have also set the linewidth property to zero
and added kernel density plots

sns.displot (x='Salary',row='Gender',data=bank,linewidth=0,kde=True);
```



```
In [ ]: # As mentioned in our discussion of R, overlaying histograms almost never
        # makes sense—the result is typically a mess, which is why SAS Enterprise
        # Guide stacks them one on top of the other (as we just did above). Here is
        # an overlaid histogram for the bank salary data. Note that the color are
        # added, so we get a third color in regions of overlap:


        sns.histplot(x='Salary', hue='Gender', data=bank, linewidth=0);
```



```
In [ ]: # A better approach is to stack kernel density plots. Note that I have
        # also added some shading to make the result look marginally cooler

        sns.kdeplot(x='Salary', hue='Gender', data=bank, fill=True);
```





## T-tests

```
In [ ]: # We use the t-test at this point to formally test the hypothesis that two
distributions have the same sample mean (and thus are “the same”–or at
least close enough). As in Excel and R, the two main preconditions to
running the test in Python are:

        # 1. Getting the data in the right format
        # 2. Determining which version of the t-test to run: equal variance or
unequal variance

female_sal=bank[bank['Gender']=="Female"]['Salary']
male_sal=bank[bank['Gender']=="Male"]['Salary']
female_sal

# Testing for equality of variance Using Levene's Test
# ensure the scipy stats module is loaded

from scipy import stats
stats.levene(female_sal,male_sal)
```

```
# from the Levene results below , we seen that the version of t-test to
run is of unequal variance since p value has e-07, meaning it is smaller
enough to treat as zero
```

```
Out[ ]: LeveneResult(statistic=26.208558688866834, pvalue=7.013920590029544e-07)
```

## Running the t-test to compare the means

```
In [ ]: import statsmodels.stats.api as sms
model=sms.CompareMeans.from_data(bank[bank['Gender']=="Female"]
['Salary'],bank[bank['Gender']=="Male"]['Salary'])
model.summary(usevar='unequal')
```

```
Out[ ]: Test for equality of means
```

	coef	std err	t	P> t	[0.025	0.975]
<b>subset #1</b>	-8.2955	2.003	-4.141	0.000	-12.283	-4.308

```
In [ ]: import statsmodels.stats.api as sms
model=sms.CompareMeans.from_data(female_sal,male_sal)
model.summary(usevar='unequal')
```

```
Out[ ]: Test for equality of means
```

	coef	std err	t	P> t	[0.025	0.975]
<b>subset #1</b>	-8.2955	2.003	-4.141	0.000	-12.283	-4.308

```
In [ ]:
```